

# Algorithms for Data Science

## Data Streams II

---

April 2, 2023

# Table of contents

Data Streams

Counting Distinct Elements

Estimating Moments in Streams

Queries over Sliding Windows

# Data Streams


Input rate is **controlled externally** – so the data processor has no control over the speed of the data

**Data streams** are:

- **infinite** – one does not know the size of the data
- **non-stationary** – the distributions of the data can change (seasonally, daily, hourly)

**Model:** infinite sequence of items  $S = (i_1, i_2, \dots, i_k \dots)$

# Stream Processing Model



**Objective:** asking **queries** on the stream – *standing* and *ad-hoc*

**Restrictions:** **storage space** and **processing time** – have to process it or we lose it forever!

# Implementation Issues

If we had enough memory / time – data streams **would be easy**

With **restrictions**:

- more efficient to get **approximate** answers
- use space-saving techniques such as **hashing**

# Problems Studied

- Sample data from a stream
- Filtering items
- Counting distinct elements
- Estimating moments
- Queries over sliding windows

# Table of contents

Data Streams

Counting Distinct Elements

Estimating Moments in Streams

Queries over Sliding Windows

# Count-Distinct Problem

**Problem:** count the number of **distinct items** in a data stream

**Applications:**

- how many different words are in webpages (spam detection)?
- how many distinct products are sold in the last week?
- how many new stars do we find in space?



# Count-Distinct Problem

**Naïve Approach** keep a set of new items found and keep a count of its size

What if we do not have enough space for all the distinct elements?

- we still want an unbiased estimator of the counts
- we accept some error in the estimation as trade-off for space

# Flajolet-Martin Approach [?]

**Algorithm** – assume we have  $N$  items in the universe:

1. pick a hash function  $h$  mapping the  $N$  items to at least  $\log_2 N$  bits
2. for each stream item  $s$ ,  $r(s)$  is the number of **trailing os in the bit representation**
  - for instance assume  $h(s) = 12$ , bit representation **1100**
  - $r(a)$  is then equal to 2
3. keep  $R = \max_s r(s)$  over the entire stream

**Estimator:** the number of distinct items seems thus far is  $2^R$ .

# Intuition on Why It Works

**Assumption:**  $h$  hashes with equal probability to all  $N$  values, the values from the stream come uniformly

$h(s)$  is a **sequence of  $\log_2 N$  bits:**

- a proportion of  $2^{-1}$  (50%) will have  $r(s) = 1$
- a proportion of  $2^{-2}$  (25%) will have  $r(s) = 2$
- **generally**, a proportion of  $2^{-r}$  will have  $r$  trailing 0s

For an **uniform hash function**, it takes thus  $1/2^{-r} = 2^r$  items before **we see one with  $r$  trailing 0s**

*Note:* it can be done with trailing 1s, or any other bit function allowing us to compute the probability

# Drawbacks and Optimization

**Main drawback:** the expectation  $E[2^R]$  can get very high

**Can fix by using multiple estimators** –  $m$  different hash functions

- taking the **average** can overestimate – if one estimator is an **outlier**
- taking the **median** is better – but it is always a power of 2
- **best approach:** hybrid, divide the hash functions in groups, compute average in each group, take the median over groups

## Minimizes space used

- only have to keep  $R$  for each hash function
- we can use as many hash functions as memory permits
- **time trade-off**: if too many – computing the hashes (and maintaining averages, medians) can be too **time** costly

# Table of contents

Data Streams

Counting Distinct Elements

Estimating Moments in Streams

Queries over Sliding Windows

# Moments of a Sequence

Assume we have a sequence/stream  $S$  having  $N$  possible distinct (ordered) values, and  $m_i$  is the number of times the  $i$ th distinct element appears in  $S$

**Moment:** the  $n$ th **moment** of a sequence  $S$  is equal to

$$\sum_{i \in S} (m_i)^n.$$

# Moments of a Sequence

Example of moments:

1. 0th moment: the number of distinct items in the stream – can be estimated using the approach presented before!
2. 1st moment: the length of the stream – easy to keep count of
3. 2nd moment: **surprise number** – how uneven the distribution is

**Challenge:** same as distinct items in stream – cannot keep all values in memory



# Surprise Number

5 distinct elements not varying much: 5 4 4 4 3

- 2nd moment (surprise number):  $1^2 + 3^2 + 1^2 = 10$

5 distinct elements with outliers: 16 1 1 1 1

- 2nd moment (surprise number):  $1^2 + 4^2 = 17$

# Alon-Matias-Szegedy Algorithm [?]

Assume a stream has a length  $n$ , and we have space to store a few variables and not all  $m_i$

We keep **some variables**  $X$ :

- $X.val$  – the value of the element
- $X.c$  – the count of that element in the stream

**Alon-Matias-Szegedy Algorithm** (AMS):

1. choose a number  $i$  between  $1$  and  $n$
2. when the stream  $S$  reaches  $i$ , set  $X.val = s_i$  and  $X.c = 1$
3. everytime the value in  $X.val$  is encountered in  $S$ , increment  $X.c$

# Using AMS for Estimating 2nd Moment

**Estimate** of the 2nd moment is:

$$n(2X_{\cdot c} - 1)$$

The estimate can be refined by using  $k$  different  $X$  variables; the estimate is then the **average** of the estimates:

$$\frac{n}{k} \sum_{i \in \{1, \dots, k\}} (2X_{i \cdot c} - 1)$$

# AMS example

Stream ( $n = 15$ ):

a b c b d a c d a b d c a a b

- **surprise number**  $5^2 + 4^2 + 3^2 + 3^2 = 59$

Keep  $X_1, X_2, X_3$ , and choose 3, 8, 13 as **random positions** in the stream:

- $X_1.val = c$ , and – at the end of the stream –  $X_1.c = 3$
- $X_2.val = d$ , and – at the end of the stream –  $X_2.c = 2$
- $X_3.val = a$ , and – at the end of the stream –  $X_3.c = 2$

The **final estimate** is:

$$15/3 \times ((2 \times 3 - 1) + (2 \times 2 - 1) + (2 \times 2 - 1)) = 55$$

# Why It Works



Let us write  $f(X) = n(2c_t - 1)$ , and  $c_t$  the number of times an item appears from time  $t$  on

We need to give a bound on the **expectation** of  $f$ :

$$\begin{aligned} \mathbb{E}[f(X)] &= \frac{1}{n} \sum_{t=1}^n n(2c_t - 1) = \sum_{i=1}^{m_i} (2i - 1) \\ &= 2 \frac{m_i(m_i + 1)}{2} - m_i = (m_i)^2 \end{aligned}$$

– **in expectation**, the formula is exactly the **second moment**!

# Estimating Higher Order Moments

The algorithm works for **any moment  $k$** , but the estimate changes

**General estimator**

$$n \left( c^k - (c - 1)^k \right)$$

*Exercise: show that it works*

## What happens when we do not know $n$ ?

- assume we can only hold  $k$  functions

## We can use Reservoir Sampling

- choose the first  $k$  times for  $k$  variables
- for  $n > k$  choose the item as a new variable with probability  $k/n$ , if chosen discard one of the previous  $k$  randomly
- in the estimator, use the current length of the stream as  $n$

# Table of contents

Data Streams

Counting Distinct Elements

Estimating Moments in Streams

Queries over Sliding Windows



**Setting:** sometimes we only need to query the last  $N$  elements of a stream – **queries over a sliding window**

- $N$  can be **very large**
- there can also be multiple stream, so keeping multiple windows is too much

*Example:* **transactions** (product was sold, ad was clicked, etc.)

# Sliding Windows

q w e r t y u i o p **a s d f g h** j k l z x c v b n m

q w e r t y u i o p a **s d f g h j** k l z x c v b n m

q w e r t y u i o p a s **d f g h j k** l z x c v b n m

q w e r t y u i o p a s d **f g h j k l** z x c v b n m

← **Past** **Future** →

J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmms.org>

# Counting Bits

**Problem:** given a stream  $S$  of 0 and 1, we want to answer queries of the form

- how many 1s are in the last  $k$  bits ( $k \leq N$ )

**Assumption:** we cannot afford to keep the most recent  $N$  bits

- but, impossible to get an exact answer without storing the entire window
- have to settle for **approximate answers**

# Non-Uniform Streams

In **uniform streams**, we can simply estimate the number of 1s by counting the number of 1s as  $a$ , 0s as  $b$  and estimate as


$$N \frac{a}{a + b}$$

**But streams are not uniform!**

# Datar-Gionis-Indyk-Motwani (DGIM) Method [?]

## Main Idea – exponential windows

- summarize regions of the streams in buckets, that are **exponentially increasing**
- keep the count for each



We can reconstruct the count of the last  $N$  bits, except we are not sure how many of the last 6 1s are included in the  $N$

J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmms.org>

# Pros and Cons

## The advantages:

- only needs  $\mathcal{O}(\log^2 N)$  bits –  $\mathcal{O}(\log N)$  counts of  $\log_2 N$  bits
- easy updates
- error in count not greater than the number of **1** in the “last” area
- if **1**s are (relatively) evenly distributed, error is no more than **50%**

## The big disadvantage:

- if all the **1**s are in the unknown area – error is unbounded!

# The DGIM Fix


**Main idea:** instead of keeping fixed sizes of buckets, keep buckets containing a fixed size of **1s**

- the windows increase exponentially – numbers of **1** kept as powers of **2**, e.g., 1 1 2 4 16

**Buckets** contain:

- the timestamp of its end – kept as timestamp modulo  **$N$** , needs  $\mathcal{O}(\log N)$  bits
- the number of **1s** in it – since powers of **2** always, it only needs  $\mathcal{O}(\log \log N)$

# Restrictions on Buckets



- at most one or two buckets of the same size
- no overlap of timestamps
- new buckets are smaller than earlier ones
- buckets are removed when end time  $> N$



# Updating Buckets

When a new item (bit) comes, drop the last bucket if end-time after  $N$

**Update** depends on the bit (0 or 1):

1. if bit is 0 – **no changes needed**
2. if bit is 1:
  - create a new bucket of size 1
  - if 3 buckets of size 1, combine oldest two in a new bucket of size 2
  - recurse on sizes

# Updating Buckets

Current state of the stream:

1001010110001011 0 10101010101011 0 101010101011 0 1010101110101 000 1011 001 0

Bit of value 1 arrives

001010110001011 0 10101010101011 0 101010101011 0 1010101110101 000 1011 001 0 1

Two orange buckets get merged into a yellow bucket

001010110001011 0 10101010101011 0 101010101011 0 1010101110101 000 1011 001 0 1

Next bit 1 arrives, new orange bucket is created, then 0 comes, then 1:

010110001011 0 10101010101011 0 101010101011 0 1010101110101 000 1011 001 0 1 0 1

Buckets get merged...

010110001011 0 10101010101011 0 101010101011 0 1010101110101 000 1011 001 0 1 0 1

State of the buckets after merging

010110001011 0 101010101010110101010101011 0 1010101110101 000 1011001 0 1 0 1

## Query:

1. sum the sizes of all buckets except the last
2. add **half the size of the last** (we do not know the proportion of the last window in  $N$ )

## Error is at most 50%:

- can be reduced by maintaining  $r$  or  $r - 1$  buckets of each size
- **error** is then at most  $\mathcal{O}(1/r)$
- **trade-off** between the number of bits and the error

# Extensions

Using  $k < N$  as a query parameter:

- want to query only the last  $k$  bits in the window  $N$
- can simply “cut” at  $k$  and use the same estimator

Sum of last  $k$  integer elements:

- integers have at most  $m$  bits
- treat each bit as a separate stream and count the **1** in last  $k$
- estimate as  $\sum_{i=0}^{m-1} c_i 2^i$  where  $c_i$  is the DGIM estimator for bit  $i$

# Acknowledgments

The contents follow Chapter 4 of [?]. Figures in slides 4, 21, 26, 29, 32, and 34 are taken from <http://www.mmds.org/>

